

High Performance Computing Lab

Class: Final Year (Computer Science and Engineering)

Year: 2022-23

PRN: 2019BTECS00089 – Piyush Pramod Mhaske

Batch: B3

Practical 9

Github link: <https://github.com/Piyush4620/2019BTECS00089HPCLab>

Hosted Link : <https://better-sidecar-c10.notion.site/HPC-038e2693a633408c8604841fc50f74e2>

Que 1 : Nbody

```
#include <math.h>
#include <stdio.h>
#include <stdlib.h>
#include "timer.h"
#include "files.h"

#define SOFTENING 1e-9f

/*
 * Each body contains x, y, and z coordinate positions,
 * as well as velocities in the x, y, and z directions.
 */

typedef struct { float x, y, z, vx, vy, vz; } Body;

/*
 * Calculate the gravitational impact of all bodies in the system
 * on all others.
 */

__global__ void bodyForce(Body *p, float dt, int N) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;

    if (tid < N) {
        float Fx = 0.0f; float Fy = 0.0f; float Fz = 0.0f;

        for (int j = 0; j < N; j++) {
            float dx = p[j].x - p[tid].x;
            float dy = p[j].y - p[tid].y;
            float dz = p[j].z - p[tid].z;
            float distSqr = dx*dx + dy*dy + dz*dz + SOFTENING;
            float invDist = rsqrtf(distSqr);
            float invDist3 = invDist * invDist * invDist;

            Fx += dx * invDist3; Fy += dy * invDist3; Fz += dz * invDist3;
```

```

    }

    p[tid].vx += dt*Fx; p[tid].vy += dt*Fy; p[tid].vz += dt*Fz;
}
}

int main(const int argc, const char** argv) {

    // The assessment will test against both 2<11 and 2<15.
    // Feel free to pass the command line argument 15 when you generate ./nbody report files
    int nBodies = 2<<11;
    if (argc > 1) nBodies = 2<<atoi(argv[1]);

    // The assessment will pass hidden initialized values to check for correctness.
    // You should not make changes to these files, or else the assessment will not work.
    const char * initialized_values;
    const char * solution_values;

    if (nBodies == 2<<11) {
        initialized_values = "09-nbody/files/initialized_4096";
        solution_values = "09-nbody/files/solution_4096";
    } else { // nBodies == 2<<15
        initialized_values = "09-nbody/files/initialized_65536";
        solution_values = "09-nbody/files/solution_65536";
    }

    if (argc > 2) initialized_values = argv[2];
    if (argc > 3) solution_values = argv[3];

    const float dt = 0.01f; // Time step
    const int nIters = 10; // Simulation iterations

    int bytes = nBodies * sizeof(Body);
    float *buf;
    cudaMallocManaged(&buf, bytes);

    Body *p = (Body*)buf;

    read_values_from_file(initialized_values, buf, bytes);

    double totalTime = 0.0;

    /*
     * This simulation will run for 10 cycles of time, calculating gravitational
     * interaction amongst bodies, and adjusting their positions to reflect.
     */

    for (int iter = 0; iter < nIters; iter++) {
        StartTimer();

        /*
         * You will likely wish to refactor the work being done in `bodyForce`,
         * and potentially the work to integrate the positions.
         */

```

```

    int threads_per_block = 128;
    int number_of_blocks = (nBodies / threads_per_block);
    bodyForce<<<number_of_blocks, threads_per_block>>>(p, dt, nBodies); // compute interbody forces
    cudaDeviceSynchronize();

    /*
    * This position integration cannot occur until this round of `bodyForce` has completed.
    * Also, the next round of `bodyForce` cannot begin until the integration is complete.
    */

    for (int i = 0 ; i < nBodies; i++) { // integrate position
        p[i].x += p[i].vx*dt;
        p[i].y += p[i].vy*dt;
        p[i].z += p[i].vz*dt;
    }

    const double tElapsed = GetTimer() / 1000.0;
    totalTime += tElapsed;
}

double avgTime = totalTime / (double)(nIters);
float billionsOfOpsPerSecond = 1e-9 * nBodies * nBodies / avgTime;
write_values_to_file(solution_values, buf, bytes);

// You will likely enjoy watching this value grow as you accelerate the application,
// but beware that a failure to correctly synchronize the device might result in
// unrealistically high values.
printf("%.3f Billion Interactions / second\n", billionsOfOpsPerSecond);

cudaFree(buf);
}

```

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
97.4	425053585	1	425053585.0	425053585	425053585	cudaMallocManaged
2.2	9599245	10	959924.5	851772	1531641	cudaDeviceSynchronize
0.3	1459656	1	1459656.0	1459656	1459656	cudaFree
0.0	212564	10	21256.4	11365	60281	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	9296107	10	929610.7	848726	1528301	bodyForce(Body*, float, int)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
53.3	238348	80	2979.4	1439	10208	[CUDA Unified Memory memcpy DtoH]
46.7	208537	15	13902.5	6816	17503	[CUDA Unified Memory memcpy HtoD]

Vector Addition:

```
#include <stdio.h>

void initWith(float num, float *a, int N)
{
    for (int i = 0; i < N; ++i)
    {
        a[i] = num;
    }
}

__global__ void addVectorsInto(float *result, float *a, float *b, int N)
{
    int start = blockIdx.x * blockDim.x + threadIdx.x;
    int stride = gridDim.x * blockDim.x;

    for (int i = start; i < N; i += stride)
    {
        result[i] = a[i] + b[i];
    }
}

void checkElementsAre(float target, float *array, int N)
{
    for (int i = 0; i < N; i++)
    {
        if (array[i] != target)
        {
            printf("FAIL: array[%d] - %0.0f does not equal %0.0f\n", i, array[i], target);
            exit(1);
        }
    }
    printf("SUCCESS! All values added correctly.\n");
}

int main()
{
    const int N = 2 << 20;
    size_t size = N * sizeof(float);

    float *a;
    float *b;
    float *c;

    cudaMallocManaged(&a, size);
    cudaMallocManaged(&b, size);
    cudaMallocManaged(&c, size);

    initWith(3, a, N);
    initWith(4, b, N);
    initWith(0, c, N);

    addVectorsInto<<<100, 1024>>>>(c, a, b, N);
    cudaDeviceSynchronize();
}
```

```

        checkElementsAre(7, c, N);

        cudaFree(a);
        cudaFree(b);
        cudaFree(c);
    }

```

Output:

```

SUCCESS! All values added correctly.

```

Profiling ;

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
95.6	246065283	3	82021761.0	16665	246016410	cudaMallocManaged
3.5	8956254	1	8956254.0	8956254	8956254	cudaDeviceSynchronize
0.9	2343818	3	781272.7	669043	879542	cudaFree
0.0	36849	1	36849.0	36849	36849	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	8950293	1	8950293.0	8950293	8950293	addVectorsInto(float*, float*, float*, int)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
79.2	5046673	591	8539.2	2142	140767	[CUDA Unified Memory memcpy HtoD]
20.8	1326491	48	27635.2	1663	159742	[CUDA Unified Memory memcpy DtoH]