

High Performance Computing Lab

Class: Final Year (Computer Science and Engineering)

Year: 2022-23

PRN: 2019BTECS00089 – Piyush Pramod Mhaske

Batch: B3

Practical 10

que 1:

1. Implement Matrix-matrix Multiplication using global memory in CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

```
#include<stdio.h>
#include<cuda.h>
#define row1 2 /* Number of rows of first matrix */
#define col1 3 /* Number of columns of first matrix */
#define row2 3 /* Number of rows of second matrix */
#define col2 2 /* Number of columns of second matrix */

__global__ void matproduct(int *l,int *m, int *n)
{
    int x=blockIdx.x;
    int y=blockIdx.y;
    int k;

    n[col2*y+x]=0;
    for(k=0;k<col1;k++)
    {
        n[col2*y+x]=n[col2*y+x]+l[col1*y+k]*m[col2*k+x];
    }
}

int main()
{
    int a[row1][col1]={1,1,1},{2,2,2};
    int b[row2][col2]={1,1},{2,2},{3,3};
    int c[row1][col2];
    int *d,*e,*f;
    int i,j;

    cudaMalloc((void **)&d,row1*col1*sizeof(int));
    cudaMalloc((void **)&e,row2*col2*sizeof(int));
    cudaMalloc((void **)&f,row1*col2*sizeof(int));

    cudaMemcpy(d,a,row1*col1*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(e,b,row2*col2*sizeof(int),cudaMemcpyHostToDevice);

    dim3 grid(col2,row1);
    /* Here we are defining two dimensional Grid(collection of blocks) structure. Syntax is dim3 grid(no. of columns,no. of rows) */

    matproduct<<<grid,1>>>(d,e,f);

    cudaMemcpy(c,f,row1*col2*sizeof(int),cudaMemcpyDeviceToHost);
    printf("\nProduct of two matrices:\n ");
    for(i=0;i<row1;i++)
    {
        for(j=0;j<col2;j++)
        {
            printf("%d\t",c[i][j]);
        }
        printf("\n");
    }

    cudaFree(d);
    cudaFree(e);
```

```

    cudaFree(f);

    return 0;
}

```

Output :

```

Product of two matrices:
    6      6
   12     12

```

Profiling :

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	272983289	3	90994429.7	3514	272974980	cudaMalloc
0.0	126573	3	42191.0	3472	116463	cudaFree
0.0	39958	3	13319.3	6282	19435	cudaMemcpy
0.0	31758	1	31758.0	31758	31758	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	4384	1	4384.0	4384	4384	matproduct(int*, int*, int*)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
59.4	2528	2	1264.0	1056	1472	[CUDA memcpy HtoD]
40.6	1728	1	1728.0	1728	1728	[CUDA memcpy DtoH]

Que 2:

2. Implement Matrix-Matrix Multiplication using shared memory in CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute

```

#include<stdio.h>
#include<cuda.h>
#define row1 2 /* Number of rows of first matrix */
#define col1 3 /* Number of columns of first matrix */
#define row2 3 /* Number of rows of second matrix */
#define col2 2 /* Number of columns of second matrix */

__global__ void matproductsharedmemory(int *l,int *m, int *n)
{
    int x=blockIdx.x;
    int y=blockIdx.y;
    __shared__ int p[col1];

    int i;
    int k=threadIdx.x;

    n[col2*y+x]=0;

    p[k]=l[col1*y+k]*m[col2*k+x];

    __syncthreads();
}

```

```

        for(i=0;i<col1;i++)
            n[col2*y+x]=n[col2*y+x]+p[i];
    }

int main()
{
    int a[row1][col1]={1,1,1},{2,2,2};
    int b[row2][col2]={1,1},{2,2},{3,3};
    int c[row1][col2];
    int *d,*e,*f;
    int i,j;

    cudaMalloc((void **)&d,row1*col1*sizeof(int));
    cudaMalloc((void **)&e,row2*col2*sizeof(int));
    cudaMalloc((void **)&f,row1*col2*sizeof(int));

    cudaMemcpy(d,a,row1*col1*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(e,b,row2*col2*sizeof(int),cudaMemcpyHostToDevice);

    dim3 grid(col2,row1);
    /* Here we are defining two dimensional Grid(collection of blocks) structure. Syntax is dim3 grid(no. of columns,no. of rows) */

    matproductsharedmemory<<<grid,col1>>>(d,e,f);

    cudaMemcpy(c,f,row1*col2*sizeof(int),cudaMemcpyDeviceToHost);

    printf("\n Product of two matrices:\n ");
    for(i=0;i<row1;i++)
    {
        for(j=0;j<col2;j++)
        {
            printf("%d\t",c[i][j]);
        }
        printf("\n");
    }

    cudaFree(d);
    cudaFree(e);
    cudaFree(f);

    return 0;
}

```

Output:

```

        Product of two matrices:
        6      6
        12     12

```

Profiling:

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	259039596	3	86346532.0	5291	259027911	cudaMalloc
0.1	129830	3	43276.7	5023	116197	cudaFree
0.0	47834	3	15944.7	6067	23793	cudaMemcpy
0.0	32437	1	32437.0	32437	32437	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	4352	1	4352.0	4352	4352	matproductsharedmemory(int*, int*, int*)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
58.5	2528	2	1264.0	1056	1472	[CUDA memcpy HtoD]
41.5	1792	1	1792.0	1792	1792	[CUDA memcpy DtoH]

Que 3:

Implement Prefix sum using CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

```
// This program computes prefix sum with warp divergence
#include <bits/stdc++.h>

using std::accumulate;
using std::generate;
using std::cout;
using std::vector;

#define SHMEM_SIZE 256

__global__ void prefixSum(int *v, int *v_r) {
    // Allocate shared memory
    __shared__ int partial_sum[SHMEM_SIZE];

    // Calculate thread ID
    int tid = blockIdx.x * blockDim.x + threadIdx.x;

    // Load elements into shared memory
    partial_sum[threadIdx.x] = v[tid];
    __syncthreads();

    // Iterate of log base 2 the block dimension
    for (int s = 1; s < blockDim.x; s *= 2) {
        // Reduce the threads performing work by half previous the previous
        // iteration each cycle
        if (threadIdx.x % (2 * s) == 0) {
            partial_sum[threadIdx.x] += partial_sum[threadIdx.x + s];
        }
        __syncthreads();
    }

    // Let the thread 0 for this block write it's result to main memory
    // Result is indexed by this block
    if (threadIdx.x == 0) {
        v_r[blockIdx.x] = partial_sum[0];
    }
}

int main() {
    // Vector size
    int N = 1 << 16;
    size_t bytes = N * sizeof(int);

    // Host data
```

```

vector<int> h_v(N);
vector<int> h_v_r(N);

// Initialize the input data
generate(begin(h_v), end(h_v), [](){ return rand() % 10; });

// Allocate device memory
int *d_v, *d_v_r;
cudaMalloc(&d_v, bytes);
cudaMalloc(&d_v_r, bytes);

// Copy to device
cudaMemcpy(d_v, h_v.data(), bytes, cudaMemcpyHostToDevice);

// TB Size
const int TB_SIZE = 256;

// Grid Size (No padding)
int GRID_SIZE = N / TB_SIZE;

// Call kernels
prefixSum<<<GRID_SIZE, TB_SIZE>>>(d_v, d_v_r);

prefixSum<<<1, TB_SIZE>>>(d_v_r, d_v_r);

// Copy to host;
cudaMemcpy(h_v_r.data(), d_v_r, bytes, cudaMemcpyDeviceToHost);

// Print the result
assert(h_v_r[0] == std::accumulate(begin(h_v), end(h_v), 0));

cout << "COMPLETED SUCCESSFULLY\n";

return 0;
}

```

Output:

In [5]: `lnvcc -std=c++14 -o saxpy 09-saxpy/01-saxpy.cu -run`

COMPLETED SUCCESSFULLY

In [6]: `lnsys profile --stats=true ./saxpy`

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.
Collecting data...
COMPLETED SUCCESSFULLY
Processing events...
Saving temporary "/tmp/nsys-report-867e-79c5-5766-32c9.qdstrm" file to disk...
Creating final output files...
Processing [=====100%]
Saved report file to "/tmp/nsys-report-867e-79c5-5766-32c9.qdrep"
Exporting 1102 events: [=====100%][5%]
]
Exported successfully to
/tmp/nsys-report-867e-79c5-5766-32c9.sqlite

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	404754150	2	202377075.0	8097	404746053	cudaMalloc
0.1	303353	2	151676.5	134307	169046	cudaMemcpy
0.0	56841	2	28420.5	11985	44856	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	23713	2	11856.5	7648	16065	prefixSum(int*, int*)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
52.0	44513	1	44513.0	44513	44513	[CUDA memcpy HtoD]
48.0	41089	1	41089.0	41089	41089	[CUDA memcpy DtoH]

CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
256.000	1	256.000	256.000	256.000	[CUDA memcpy DtoH]
256.000	1	256.000	256.000	256.000	[CUDA memcpy HtoD]

Que 4:

Implement 2D Convolution using shared memory using CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

```
#include "cuda_runtime.h"
#include "device_launch_parameters.h"

#include <stdio.h>
#include <cstdlib>
#include <time.h>

#define BLOCK_SIZE 32
#define WA 512
#define HA 512
#define HC 3
#define WC 3
#define WB (WA - WC + 1)
#define HB (HA - HC + 1)

__global__ void Convolution(float* A, float* B, float* C, int numRows, int numACols, int numBRows, int numBCols, int numCRows, int numCCols)
{
    int col = blockIdx.x * (BLOCK_SIZE - WC + 1) + threadIdx.x;
    int row = blockIdx.y * (BLOCK_SIZE - WC + 1) + threadIdx.y;
    int row_i = row - WC + 1;
    int col_i = col - WC + 1;

    float tmp = 0;

    __shared__ float shm[BLOCK_SIZE][BLOCK_SIZE];

    if (row_i < WA && row_i >= 0 && col_i < WA && col_i >= 0)
    {
        shm[threadIdx.y][threadIdx.x] = A[col_i * WA + row_i];
    }
    else
    {
        shm[threadIdx.y][threadIdx.x] = 0;
    }

    __syncthreads();

    if (threadIdx.y < (BLOCK_SIZE - WC + 1) && threadIdx.x < (BLOCK_SIZE - WC + 1) && row < (WB - WC + 1) && col < (WB - WC + 1))
    {
        for (int i = 0; i < WC; i++)
            for (int j = 0; j < WC; j++)
                tmp += shm[threadIdx.y + i][threadIdx.x + j] * C[j*WC + i];
        B[col*WB + row] = tmp;
    }
}

void randomInit(float* data, int size)
{

```

```

    for (int i = 0; i < size; ++i)
        data[i] = rand() / (float)RAND_MAX;
}

int main(int argc, char** argv)
{
    srand(2006);
    cudaError_t error;
    cudaEvent_t start_G, stop_G;

    cudaEventCreate(&start_G);
    cudaEventCreate(&stop_G);

    unsigned int size_A = WA * HA;
    unsigned int mem_size_A = sizeof(float) * size_A;
    float* h_A = (float*)malloc(mem_size_A);

    unsigned int size_B = WB * HB;
    unsigned int mem_size_B = sizeof(float) * size_B;
    float* h_B = (float*)malloc(mem_size_B);

    unsigned int size_C = WC * HC;
    unsigned int mem_size_C = sizeof(float) * size_C;
    float* h_C = (float*)malloc(mem_size_C);

    randomInit(h_A, size_A);
    randomInit(h_C, size_C);

    float* d_A;
    float* d_B;
    float* d_C;

    cudaMalloc((void**)&d_A, mem_size_A);

    cudaMalloc((void**)&d_B, mem_size_B);

    cudaMalloc((void**)&d_C, mem_size_C);

    cudaMemcpy(d_A, h_A, mem_size_A, cudaMemcpyHostToDevice);

    cudaMemcpy(d_C, h_C, mem_size_C, cudaMemcpyHostToDevice);

    dim3 threads(BLOCK_SIZE, BLOCK_SIZE);
    dim3 grid((WB - 1) / (BLOCK_SIZE - WC + 1), (WB - 1) / (BLOCK_SIZE - WC + 1));

    Convolution << < grid, threads >> >(d_A, d_B, d_C, HA, WA, HB, WB, HC, WC);

    cudaEventRecord(start_G);

    Convolution << < grid, threads >> >(d_A, d_B, d_C, HA, WA, HB, WB, HC, WC);

    cudaDeviceSynchronize();

    cudaEventRecord(stop_G);

    cudaEventSynchronize(stop_G);

    cudaMemcpy(h_B, d_B, mem_size_B, cudaMemcpyDeviceToHost);

    float miliseconds = 0;
    cudaEventElapsedTime(&miliseconds, start_G, stop_G);

    printf("Time took to compute matrix A of dimensions %d x %d on GPU is %f ms \n \n", WA, HA, miliseconds);

    // for (int i = 0; i < HB; i++)
    // {
    //     for (int j = 0; j < WB; j++)
    //     {
    //         printf("%f ", h_B[i*HB + j]);
    //     }
    //     printf("\n");
    // }

    free(h_A);
    free(h_B);
    free(h_C);
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);
}

```


```

return EXIT_SUCCESS;
}

```

Output :

For different Values we get this trend:



Statistics:

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	2224422293	2	1112211146.5	1017	2224421276	cudaEventCreate
0.0	773053	3	257684.3	175336	374485	cudaMemcpy
0.0	743391	1	743391.0	743391	743391	cudaEventSynchronize
0.0	289421	3	96473.7	4507	279814	cudaMalloc
0.0	160776	1	160776.0	160776	160776	cudaDeviceSynchronize
0.0	150783	3	50261.0	4046	136878	cudaFree
0.0	66721	2	33360.5	6589	60132	cudaLaunchKernel
0.0	7025	2	3512.5	1804	5221	cudaEventRecord

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	167838	2	83919.0	83807	84031	Convolution(float*, float*, float*, int, int, int, int, int, int)
48.2	158080	1	158080.0	158080	158080	[CUDA memcpy DtoH]

CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
1024.035	2	512.018	0.035	1024.000	[CUDA memcpy HtoD]
1016.016	1	1016.016	1016.016	1016.016	[CUDA memcpy DtoH]